

# Arquitectura de Computadores

## **Capítulo 4. Técnicas Avanzadas en Paralelismo. Parte 1.**

**Based on the original material of the book:  
D.A. Patterson y J.L. Hennessy “Computer Organization  
and Design: The Hardware/Software Interface” 4<sup>th</sup>  
edition.**

**Escuela Politécnica Superior  
Universidad Autónoma de Madrid**


**Profesores:  
G130 y 131: Gustavo Sutter  
G135 y 139: Iván González Martínez  
G132: Francisco Javier Gómez Arribas**

# Processor Architectures


- Scalar processor : Pipelining, ILP
- Superscalar processor:
  - Multiple issue
  - Dynamic scheduling
- VLIW processor:
  - Static multiple issue
- MultiThreading
- Simultaneous Multithreading
- Multicore

# Increasing the Processor Performance


- Scalar
- Pipeline
- Superscalar
- Very Long Instruction Word (VLIW)




Scalar



Pipelined



Superscalar



VLIW ("EPIC")

# Instruction-Level Parallelism (ILP)

- Pipelining: executing multiple instructions in parallel
- To increase ILP
  - Deeper pipeline
    - Less work per stage  $\Rightarrow$  shorter clock cycle
  - Multiple issue
    - Replicate pipeline stages  $\Rightarrow$  multiple pipelines
    - Start multiple instructions per clock cycle
    - $CPI < 1$ , so use Instructions Per Cycle (IPC)
    - E.g., 4GHz 4-way multiple-issue
      - 16 BIPS, peak  $CPI = 0.25$ , peak  $IPC = 4$
    - But dependencies reduce this in practice

# Multiple Issue

- Static multiple issue
  - Compiler groups instructions to be issued together
  - Packages them into “issue slots”
  - Compiler detects and avoids hazards
- Dynamic multiple issue
  - CPU examines instruction stream and chooses instructions to issue each cycle
  - Compiler can help by reordering instructions
  - CPU resolves hazards using advanced techniques at runtime

# Static Multiple Issue

- Compiler groups instructions into “issue packets”
  - Group of instructions that can be issued on a single cycle
  - Determined by pipeline resources required
- Think of an issue packet as a very long instruction
  - Specifies multiple concurrent operations
  - $\Rightarrow$  Very Long Instruction Word (VLIW)

# Scheduling Static Multiple Issue

- Compiler must remove some/all hazards
  - Reorder instructions into issue packets
  - No dependencies within a packet
  - Possibly some dependencies between packets
    - Varies between ISAs; compiler must know!
  - Pad with nop if necessary

# Dynamic Multiple Issue

- “Superscalar” processors
- CPU decides whether to issue 0, 1, 2, ... each cycle
  - Avoiding structural and data hazards
- Avoids the need for compiler scheduling
  - Though it may still help
  - Code semantics ensured by the CPU



# Dynamic Pipeline Scheduling

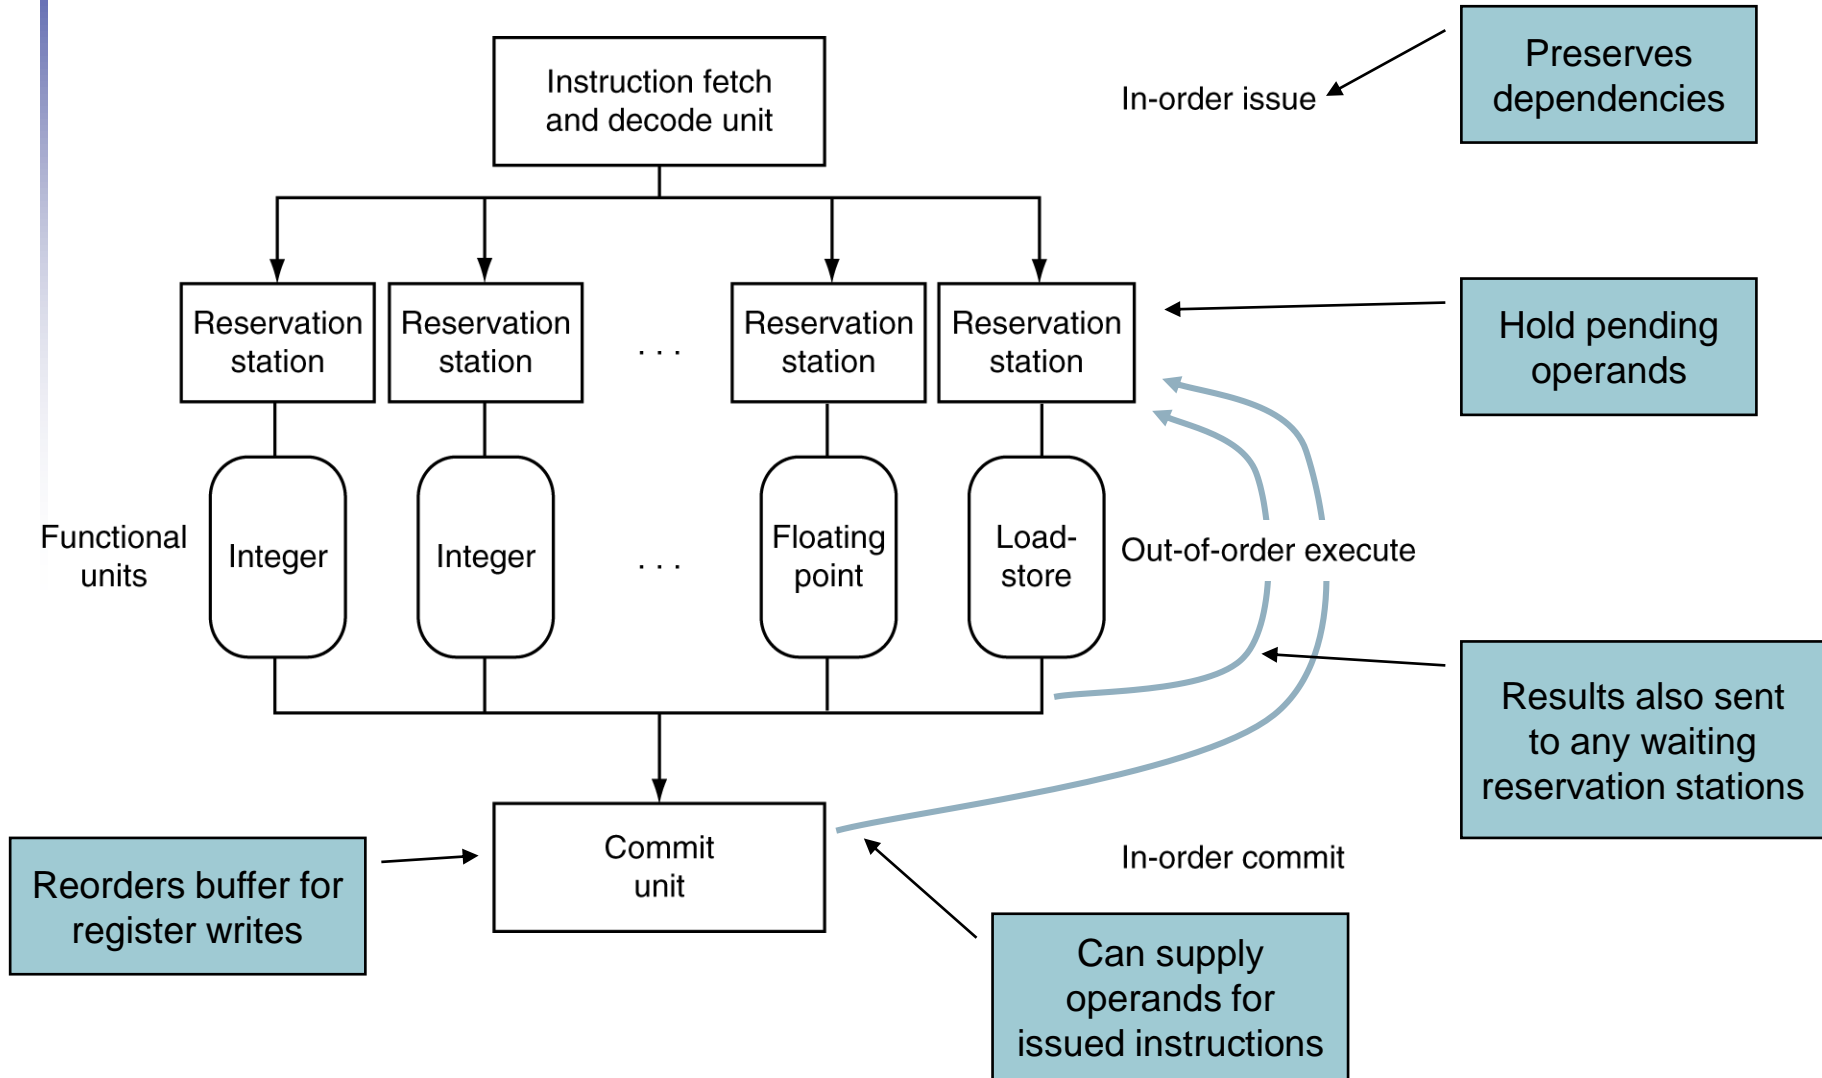
- Allow the CPU to execute instructions out of order to avoid stalls
  - But commit result to registers in order

- Example

```
lw      $t0, 20($s2)
addu    $t1, $t0, $t2
sub      $s4, $s4, $t3
slli    $t5, $s4, 20
```

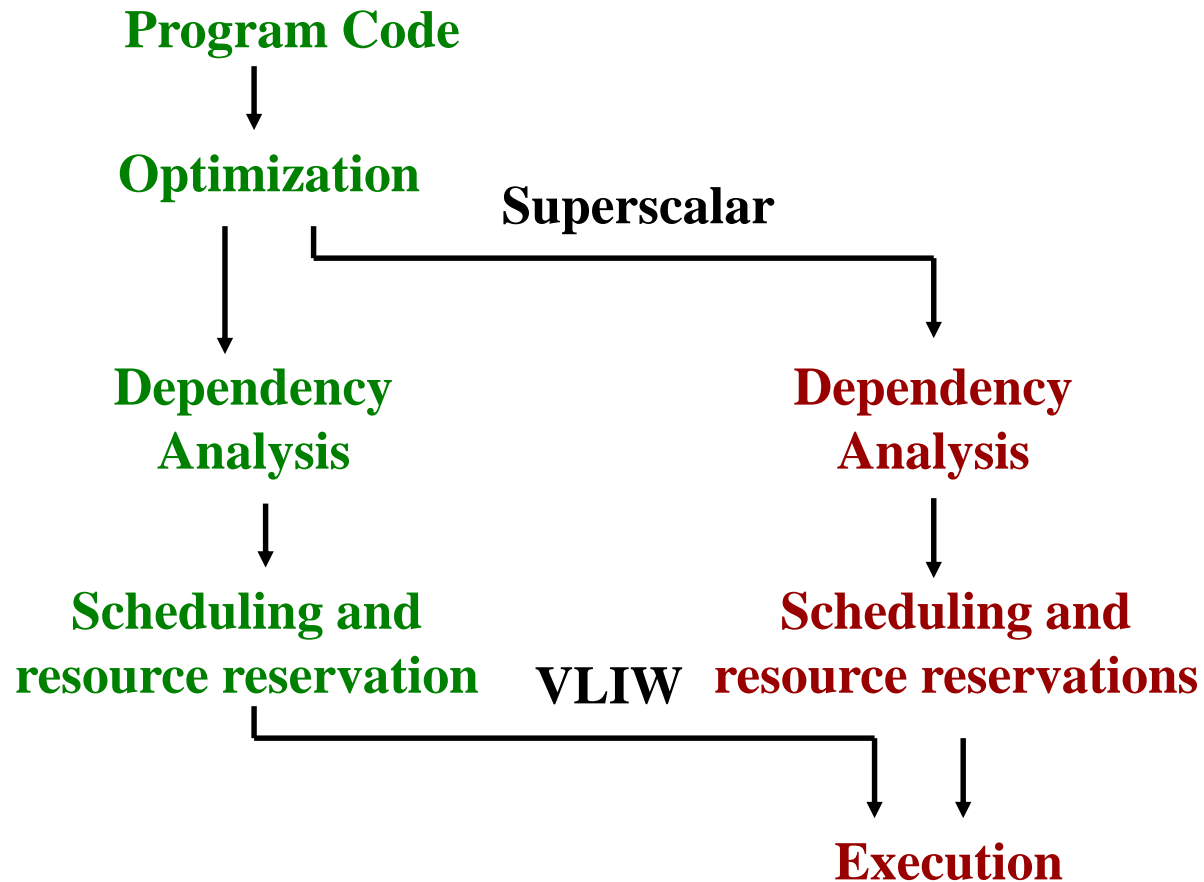
- Can start sub while addu is waiting for lw

# Dynamically Scheduled CPU




# Superscalar vs VLIW

**Different division of tasks between the SOFTWARE & HARDWARE**

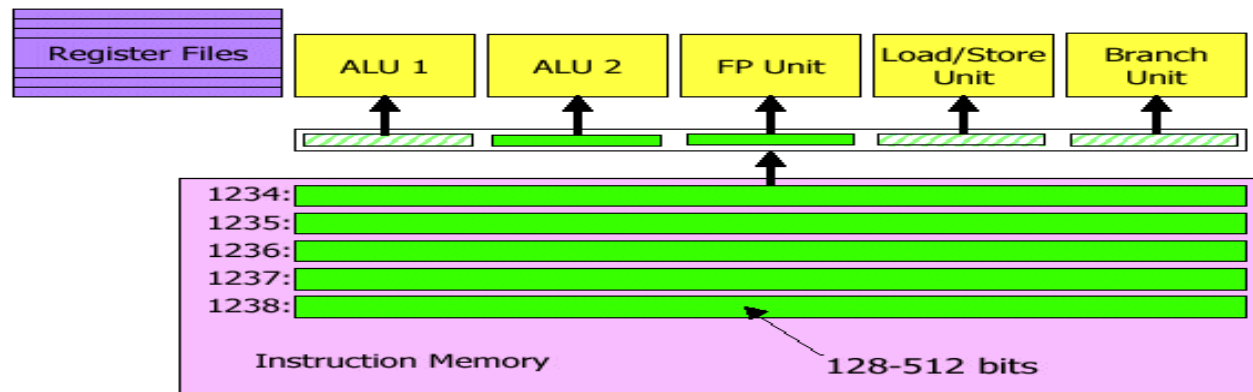


# Superscalar vs VLIW

## Superscalar



## VLIW



# Why Do Dynamic Scheduling?

- Why not just let the compiler schedule code?
- Not all stalls are predictable
  - e.g., cache misses
- Can't always schedule around branches
  - Branch outcome is dynamically determined
- Different implementations of an ISA have different latencies and hazards

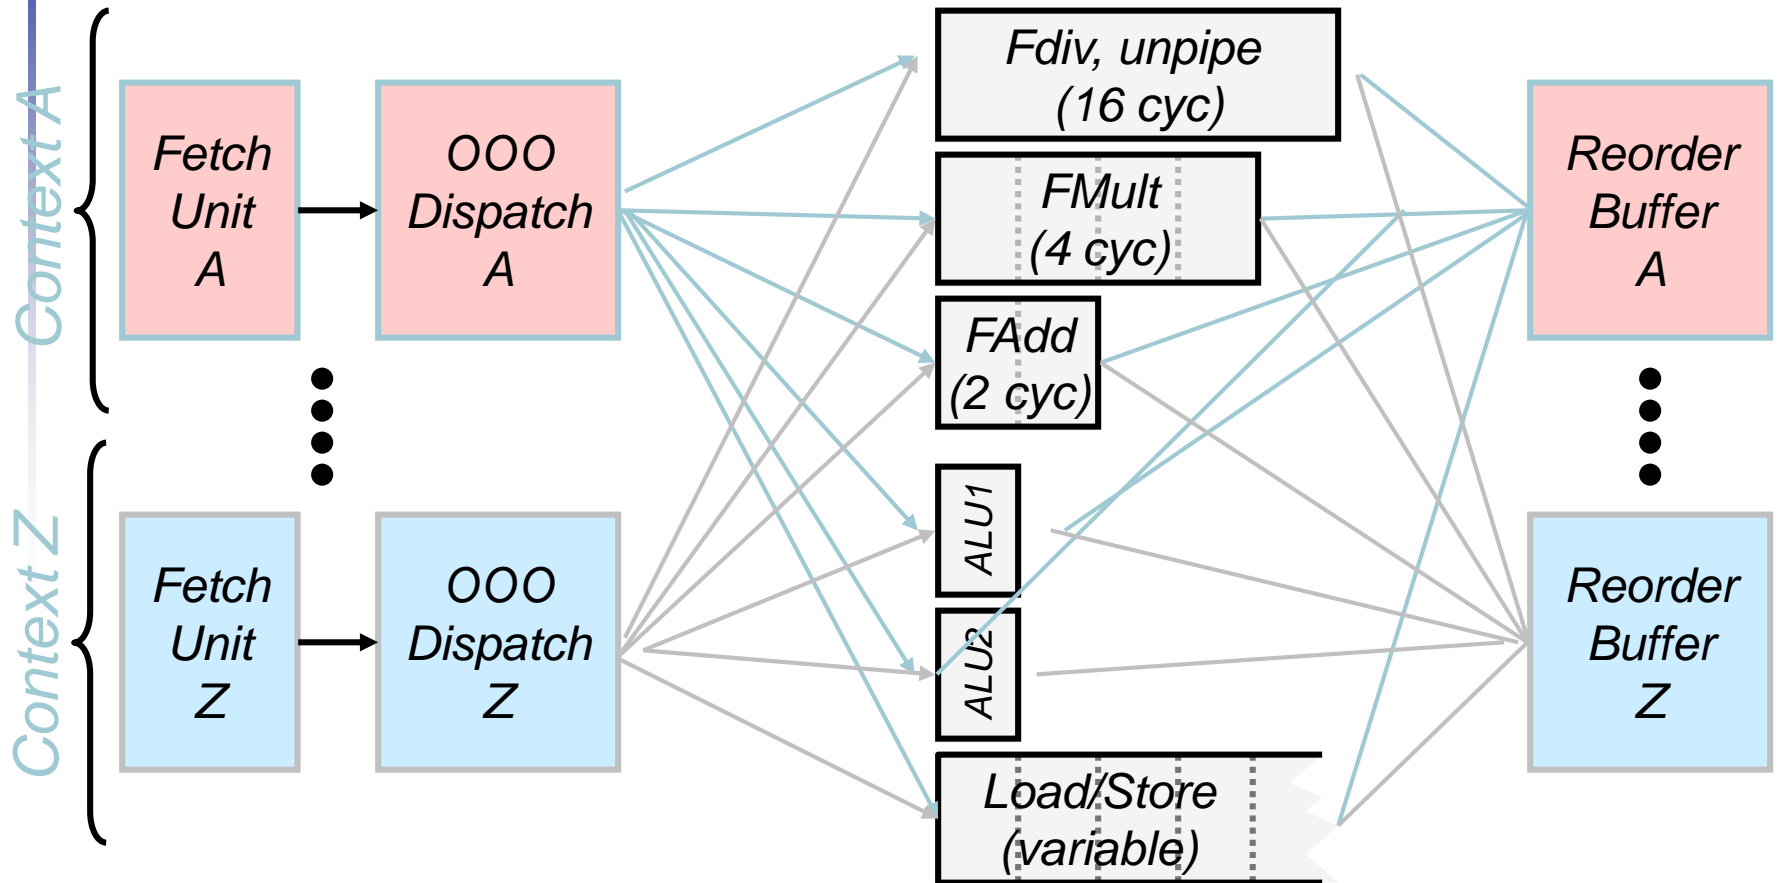
# Does Multiple Issue Work?

- Yes, but not as much as we'd like
- Programs have real dependencies that limit ILP
- Some dependencies are hard to eliminate
  - e.g., pointer aliasing
- Some parallelism is hard to expose
  - Limited window size during instruction issue
- Memory delays and limited bandwidth
  - Hard to keep pipelines full
- Speculation can help if done well

# Multithreading

- Performing multiple threads of execution in parallel
  - Replicate registers, PC, etc.
  - Fast switching between threads
- Fine-grain multithreading
  - Switch threads after each cycle
  - Interleave instruction execution
  - If one thread stalls, others are executed
- Coarse-grain multithreading
  - Only switch on long stall (e.g., L2-cache miss)
  - Simplifies hardware, but doesn't hide short stalls (eg, data hazards)

# Simultaneous Multi-Threading [Eggers, et al.]



■ Se comparten dinámicamente unidades funcionales entre múltiple threads

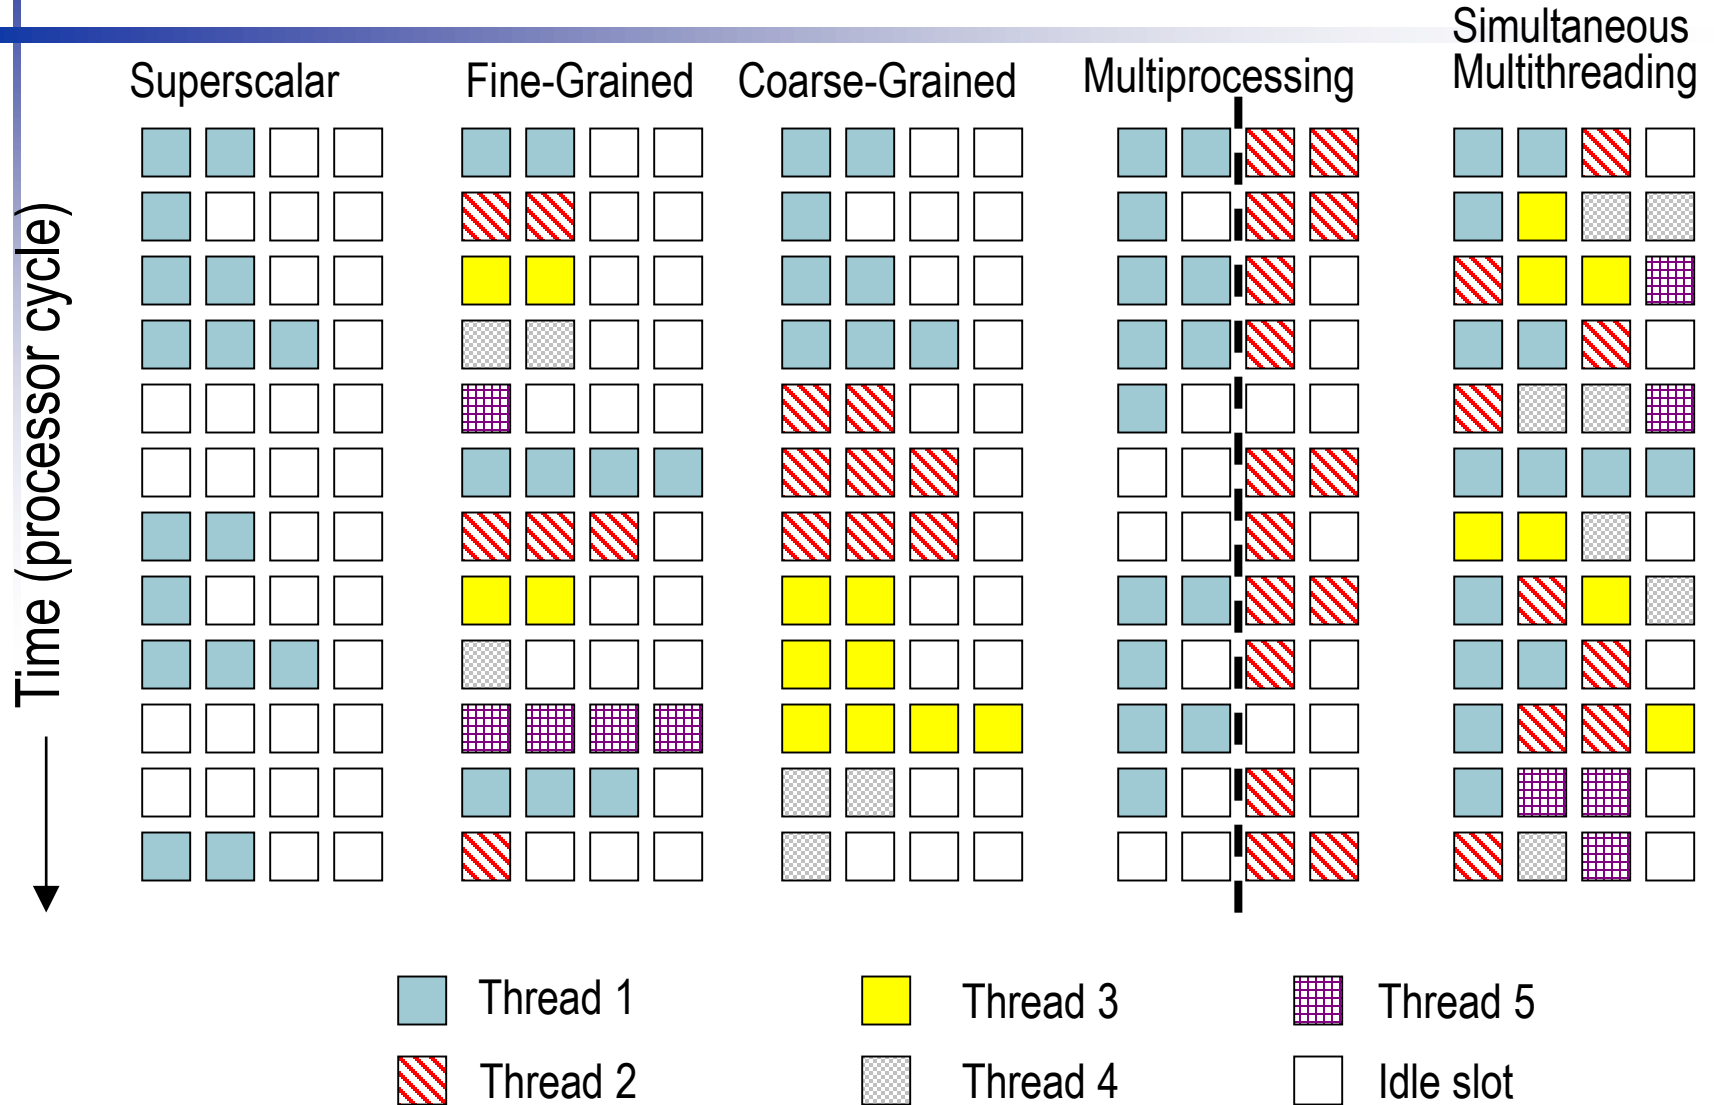
⇒ *mayor utilización* ⇒ *mayor rendimiento*





# Simultaneous Multithreading

- In multiple-issue dynamically scheduled processor
  - Schedule instructions from multiple threads
  - Instructions from independent threads execute when function units are available
  - Within threads, dependencies handled by scheduling and register renaming
- Example: Intel Pentium-4 HT
  - Two threads: duplicated registers, shared function units and caches

# TLP: Thread Level Parallelism



# Multicore Processors

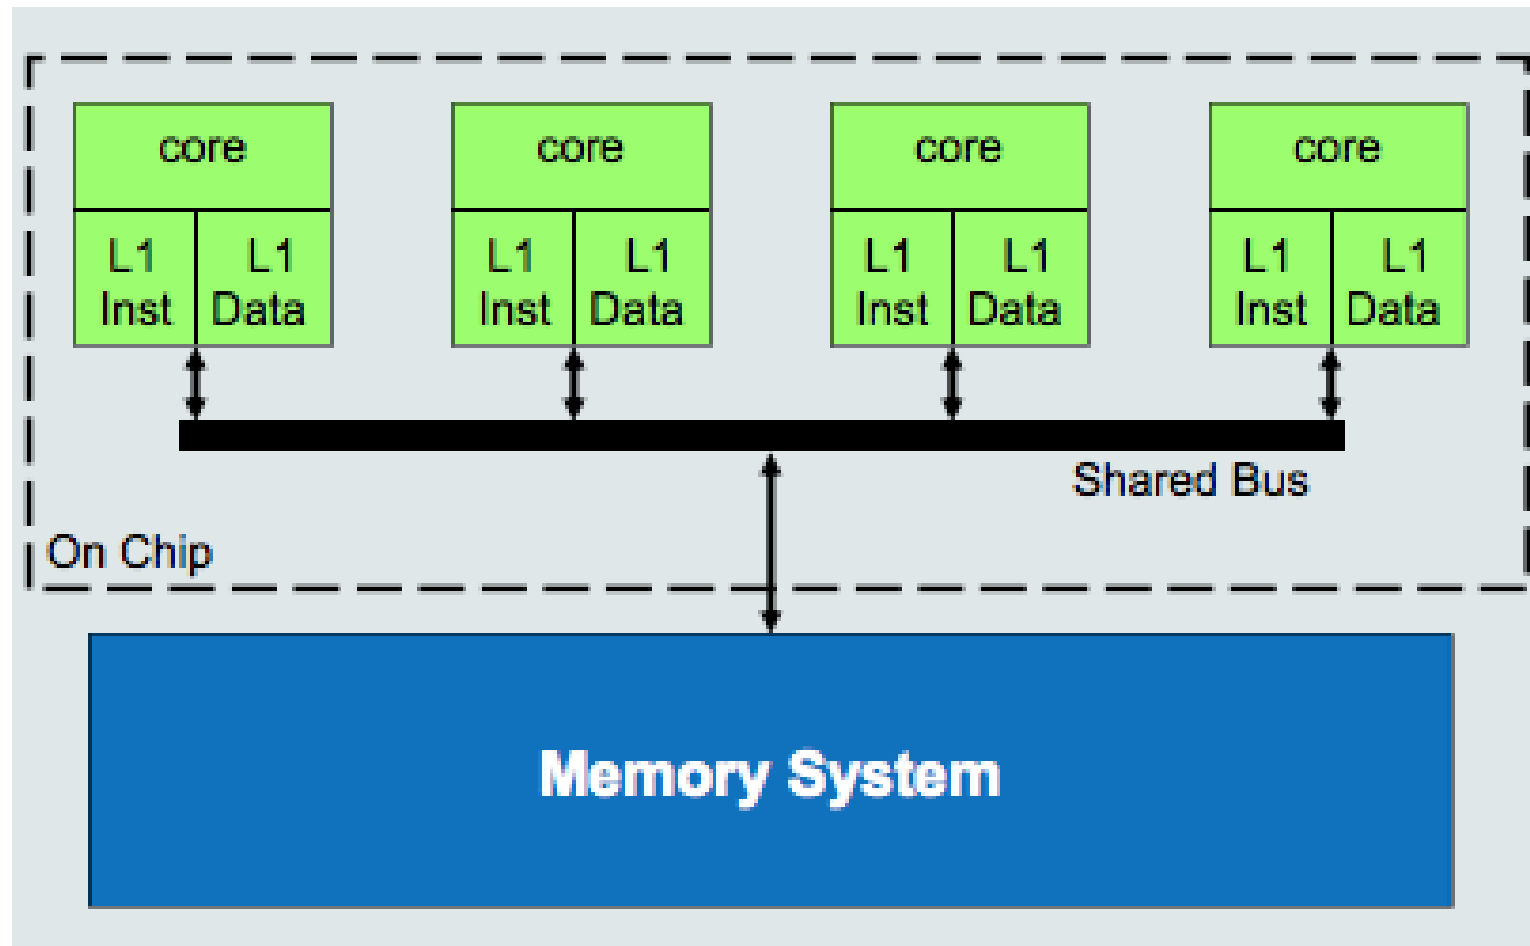


- Multi-core, each core Multi-threaded
- Shared cache and front side bus
- Each core has different Vdd & Freq


# Multi-Core

- Limitations of single core architecture:
  - High power consumption due to high clock rates (usually, ca. 2-3% power increase per 1% performance increase).
  - Heat generation (cooling is expensive).
  - Limited parallelism (Instruction-Level Parallelism only).
  - Design time and complexity increased due to complex methods to increase ILP.
- Many new applications are multi-threaded, suitable for multi-core.
  - Ex. Multimedia applications.
- Integration of multiple processor cores on a single chip


# A simple Multi-Core model




# Alternative models




(a) Dedicated L1 cache  
ARM11 MP-Core



(b) Dedicated L2 cache  
AMD Opteron

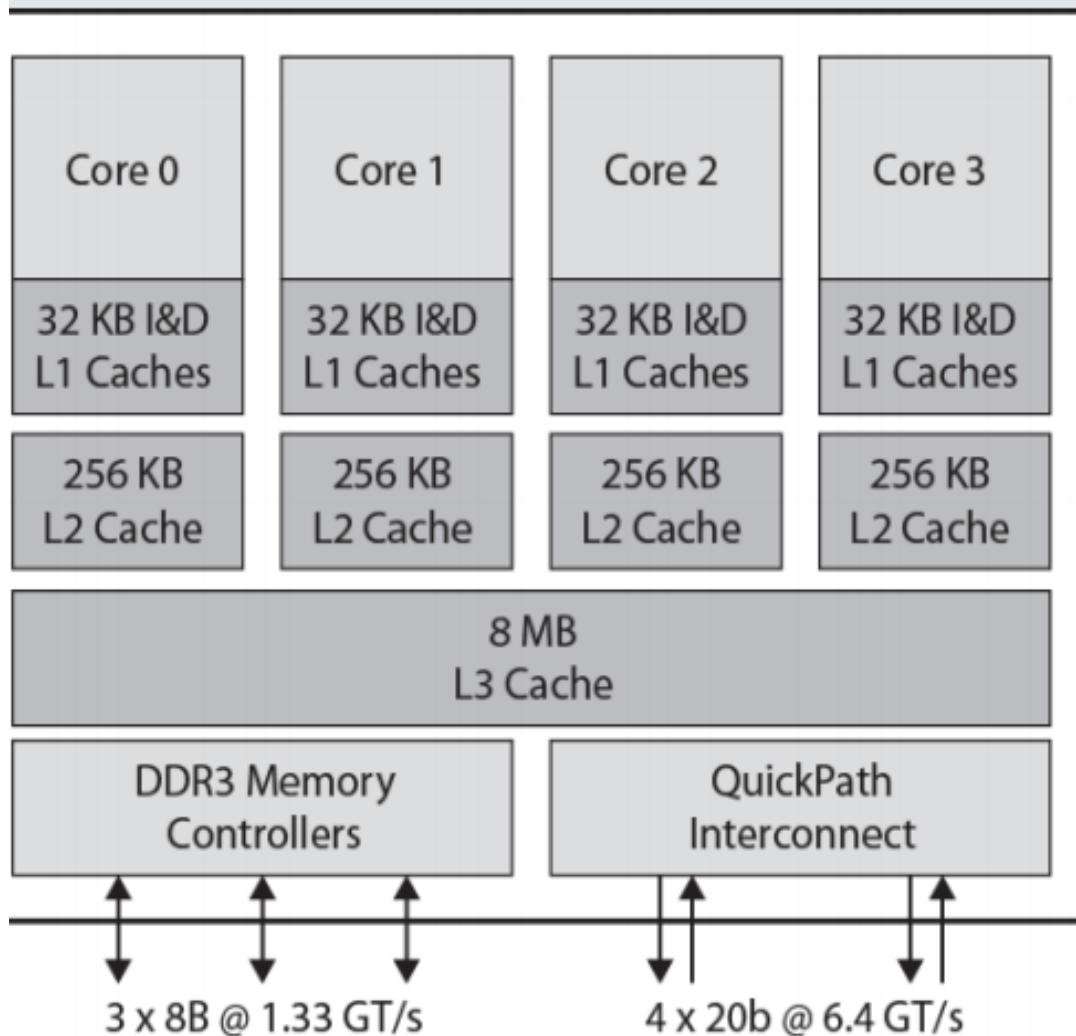


(c) Shared L2 cache  
Intel Core Duo



(d) Shared L3 cache  
Intel Core i7 and AMD Ryzen 7

# Intel Core i7



- Four (or six) identical x86 processors.
- Each with its own L1 split, and L2 unified caches.
- Shared on-chip L3 cache to speed up inter-processor communication
- High speed link between processor chips.
- Up to 4 GHz clock frequency.

# Programming for Multi-Core

There must be many threads or processes:

- Multiple applications running on the same machine.
  - Multi-tasking is becoming very common.
  - OS software tends to run many threads as a part of its normal operation.
- OS scheduler should map threads to different cores:
  - To balance the work load; or
  - To avoid hot spots due to heat generation.